

# WHEN IS MODEL SOUPING TASTY? SIMILARITY, TRANSITIVITY, AND ROBUSTNESS

000  
001  
002  
003  
004  
005 **Anonymous authors**  
006 Paper under double-blind review  
007  
008  
009  
010  
011  
012  
013  
014  
015  
016  
017  
018  
019  
020  
021  
022  
023  
024  
025  
026  
027  
028  
029  
030  
031  
032  
033  
034  
035  
036  
037  
038  
039  
040  
041  
042  
043  
044  
045  
046  
047  
048  
049  
050  
051  
052  
053  
**ABSTRACT**

Model souping is a technique in which the parameters of models are averaged, often leading to improved performance over constituent models without increasing inference cost. However, the specific conditions required for success are not well understood, particularly regarding the trade-off between model diversity and stability. We analyse over 5,000 binary ResNet-50 soups trained on CIFAR-100, with diversity controlled by branching ingredients from a shared training trajectory at varying epochs. We find that effective souping requires a balance: models must be similar enough to avoid model collapse, but diverse enough to yield improvements. Furthermore, we provide empirical evidence for the hypothesis that souping works by averaging within a low-loss basin. We also observe that soup gains on corrupted data are strongly correlated with those on in-distribution data. Finally, we compare souping to other parameter averaging methods and find that XXX. Code and experiments are available at: <https://anonymous.4open.science/r/too-salty-478E/>.

## 1 INTRODUCTION

Wortsman et al. (2022) introduced the idea of *souping* by showing that averaging model parameters produced by different fine-tuning trajectories often leads to better generalization than any individual ingredient. This not only captures the benefit of multiple adaptation paths but also introduces an additional adaptive mechanism, as Croce et al. (2023) suggest that dynamically adjusting soup weights enables intermediate behaviours that can better match a range of distributional shifts. In contrast to ensembling which combines the *outputs* of models, souping does not increase the computational cost of inference, requiring only a single forward pass.

Wortsman et al. (2022) hypothesise that souping works because fine-tuned models often lie in the same low-loss basin. The convex combination of their weights is expected to remain in the low-loss basin while reducing variance introduced by a noisy training procedure. They find that when the angle formed by the pre-trained model and the two models to be souped is larger, there is a greater performance boost from souping. A wider angle suggests that fine-tuned trajectories are more diverse, thus more variance is reduced by averaging. Understanding why such averaging remains in low-loss regions, how diversity among fine-tuned models contributes to robustness, and when souping is beneficial is therefore essential for studying adaptation more broadly.

### 1.1 RELATED WORK

**Souping:** Souping has been used in a variety of settings. Croce et al. (2023) soup models trained to be robust to different distribution shifts and Ramé et al. (2023) soup ingredients trained on different tasks. Both cases lead to better generalisation. Jang et al. (2025) further explore the optimal soup between just two fine-tuned models by considering the angle formed by the two training trajectories in a layer-wise fashion.

**Stochastic Weight Averaging (SWA):** Souping is related to the idea of SWA Izmailov et al. (2019). In SWA, the ingredients of the soup come from different steps along the same training trajectory. By contrast, souping averages models from independent training trajectories. **Exponential Moving Average (EMA):** Another method for averaging weights over training is EMA Morales-Brotóns et al.

(2024). Here, the weights at each training step are combined with the previous average using an exponential decay.

**Stability Analysis:** Souping success requires models to be ‘compatible’ in the sense that averaging their weights has low loss. Frankle et al. (2020) define *stability to SGD noise*, whereby at some point during training, models become robust to the noise from SGD in the sense that all possible minima obtained by training from that point onwards lie in the same low-loss basin.

## 1.2 OUR CONTRIBUTIONS

Following these works, we seek to better understand souping by conducting a series of experiments addressing the following questions:

**How much shared training is required for souping to be effective?** We investigate varying the number of shared pre-training epochs before splitting into fine-tuned variants.

**Can we predict the effectiveness of souping using similarity measures?** If two models are very similar, it may be more likely that they can soup effectively.

**Is souping transitive —** If model  $A$  soups with  $B$ , and  $B$  soups with  $C$ , will  $A$  soup with  $C$ ?

**Does souping in-distribution predict souping out-of-distribution?** While souping has been shown to help with robustness to distribution shifts, we seek to answer how correlated the soup gains are between in-distribution and out-of-distribution data.

Further experiments investigating the effect of permuting models Ainsworth et al. (2023) prior to souping and how souping affects robustness can be found in Appendix A.1 and A9 respectively.

## 1.3 SOUPS AND SOUP GAIN

For this work, we only soup pairs of models  $\theta_A$  and  $\theta_B$  using the simple arithmetic mean. That is,  $\theta_{\text{soup}} = \frac{1}{2}\theta_A + \frac{1}{2}\theta_B$ . We do not consider any other weighted average in order to save on computational cost. Work by Ainsworth et al. (2023), shows many of the loss barriers they find have the most extreme behaviour at the midpoint. Thus we assume that the midpoint serves as an accurate summary statistic of souping performance over all possible convex combinations.

We define the *soup gain* of a pair of models  $\theta_A, \theta_B$  as the decrease in loss over a test set obtained by souping the models. That is, soup gain =  $\min\{L(\theta_A), L(\theta_B)\} - L(\theta_{\text{soup}})$  where  $L(\theta)$  denotes the test loss obtained using model  $\theta$ . Soup gain can also be computed in terms of accuracy rather than loss. We use the soup gain as the primary measure of the effectiveness of souping in our experiments. We prefer to compare to the minimum rather than the mean of the ingredients as the purpose of a soup should be to improve over its ingredients.

## 2 EXPERIMENTS

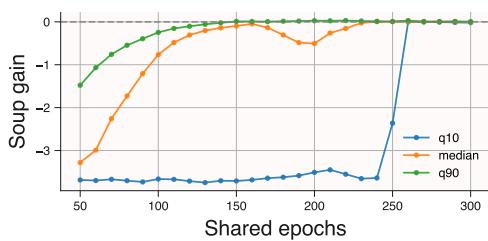
### 2.1 METHOD

We train a baseline model for image classification on the CIFAR-100 dataset with ResNet-50 He et al. (2016) using a well-known set of hyper-parameters Dadalto (2023) with image reflection and random translation with padding for data augmentation. We hold out 5% of the 50,000 training images as a validation set, as well as 10,000 images for testing. The baseline model is saved every 10 epochs. From each save point, we branch off and train 4 new models with different optimizer settings. For details, see Appendix Tables 1 and 2. All models are each trained to convergence, with the best validation scored model weights saved for the experiments. This process is illustrated in Figure A2. A total of 4 variants and 26 branch points were trained, yielding 104 related models and 5,356 binary souping combinations for analysis.

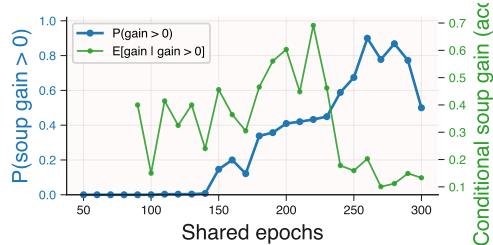
### 2.2 SHARED EPOCHS AND SOUP GAIN

We group the soups by the number of shared epochs between the two ingredients before they diverge into different training trajectories. For example, a pair of models branched form the baseline at epoch 50 and 100 respectively share 50 epochs. In Figure 1, we plot quantiles of soup gain as a function of shared epochs. We observe the distribution of soup gains shifts positively as the number of shared

108 epochs increases. However, the soup gain is often large and negative until around epoch 150. After  
 109 epoch 250, nearly all soups are approximately neutral.  
 110



120 Figure 1: Quantiles of soup gain vs shared  
 121 epochs.

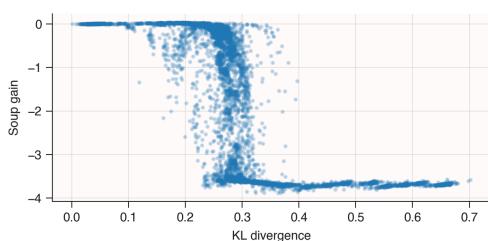


122 Figure 2: Probability of positive soup gain and  
 123 conditional expected gain vs shared epochs.

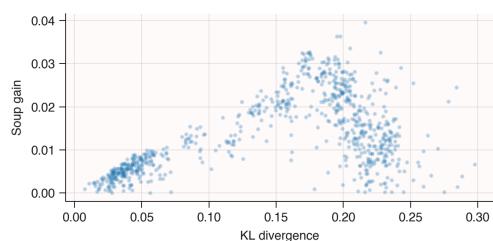
124 To better illustrate how much souping helps, we plot the probability of positive soup gain and the  
 125 average soup gain in accuracy for such positive soups in Figure 2. The probability of positive soup  
 126 gain increases with shared epochs, reaching around 80% after 250 shared epochs. The conditional  
 127 expected soup gain is noisy, at around 0.5% accuracy improvement when souping works. Towards  
 128 the upper end of shared epochs, the gains decrease to around 0.2%. This suggests that while souping  
 129 becomes more likely to work with increased shared training, the magnitude of gain decreases when  
 130 the models are too similar.

### 131 2.3 PREDICTING SOUPABILITY WITH SIMILARITY

132 Given two models, can we predict whether or not they will soup? To test this hypothesis, we compute  
 133 a variety of similarity and distance metrics between pairs of models. We find that all metrics perform  
 134 similarly. A plot for all metrics can be found in Figure A3. We arbitrarily choose to show the KL  
 135 divergence between the outputs of the ingredients in Figure 3 as an example. There is a strong  
 136 correlation between KL divergence and soup gain with a Spearman correlation of -0.86.  
 137 However, many of these soups have very poor performance. Figure 4 shows the soup gain against KL  
 138 divergence for only those models with positive soup gain, with a moderate positive correlation with a  
 139 Spearman correlation of 0.39. Models must be sufficiently similar in order to soup, but to be effective,  
 140 the models must also be sufficiently different. Balancing these two effects is key to tasty soups.



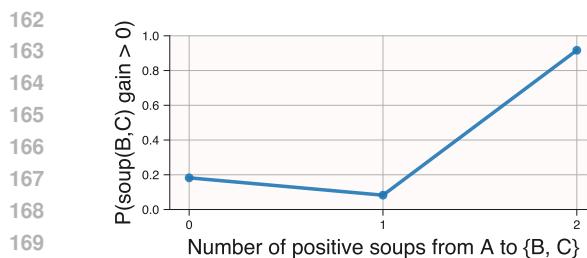
141 Figure 3: KL vs soup gain (Spearman -0.86).



142 Figure 4: KL vs soup gain, positive soups only  
 143 (Spearman 0.39).

### 144 2.4 IS SOUPING TRANSITIVE?

145 If model  $A$  soups successfully with models  $B$  and  $C$ , will  $B$  and  $C$  also soup successfully? If models  
 146 soups together when they lie in the same low loss region, then transitivity should hold. To test this  
 147 hypothesis, we consider all triplets of models  $(A, B, C)$ . We plot the probability that  $B$  and  $C$  soup  
 148 against the number of positive soups involving  $A$  in Figure 5. We observe that the probability that  $B$   
 149 and  $C$  soup is only high when  $A$  soups with both  $B$  and  $C$ . We also plot the soup gain between  $B$   
 150 and  $C$  against the minimum soup gain of  $(A, B)$  and  $(A, C)$  in Figure A5, noting a moderate positive  
 151 Spearman correlation of 0.64.



171 Figure 5: Probability of positive soup gain of  $B$   
172 and  $C$  vs positive soups with  $A$ .  
173

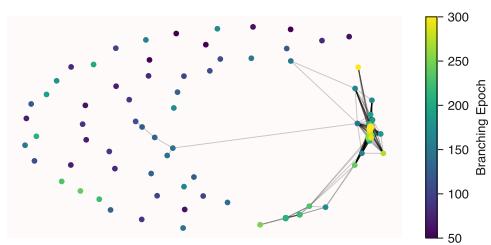


Figure 6: 2D embedding of 104 models using soup-gain distance; edges indicate positive soups.

Following transitivity, we investigate whether there are clusters of models that soup well together. We embed our 104 models into 2D using soup gain as a distance metric and plot them in Figure 6. The details of generating this plot can be found in Section A.4. We find that most successful soups lie in a single cluster of models all branched from later on in the training procedure. We conclude that souping is approximately transitive. There are many counter-examples, but this work supports the idea that, in general, soups lie in a single loss basin.

## 2.5 SOUPING FOR ROBUSTNESS TO CORRUPTION

All experiments thus far have measured soup gain on a held-out test set. However, souping has also been used for robustness to distribution shift (Croce et al., 2023). To establish whether souping for in-distribution performance also increases out-of-distribution performance, we compute the soup gain on CIFAR-100C (Hendrycks & Dietterich, 2019) with severity level 3. The soup gains on test and corrupted data correlate very well, with a Spearman correlation of 0.99. A scatterplot can be found in Figure A6. The positive trend still holds when conditioning on only soups with positive soup gain on the test set, with a Pearson correlation of 0.61. A plot can be found in Figure A8. In-distribution performance improvement transfers to unseen target distributions.

We also plot the probability of positive soup gain on corrupted data as a function of shared epochs in Figure A7. The probability of positive gain increases with the number of shared epochs, for both clean and corrupted data. However, the corrupted data consistently has a slightly lower probability of positive gain.

## 3 CONCLUSION

In this work, we have conducted a series of experiments to further our understanding of souping. We find that ingredients must have sufficiently many shared epochs of training in order to be compatible but that too many shared epochs and the soup gain is minimal. We also find that various similarity measures between models correlate similarly with soup gain. Similar ingredients are less likely to collapse when souped, but very similar ingredients yield smaller soup gains. Thus, the right balance must be struck for the most effective souping. When souping, we encourage practitioners to test a range of similarities of ingredients to ensure they are finding an optimal soup. Our experiments showing that souping is mostly transitive support the low-loss basin hypothesis. Finally, we find that soup gains on in-distribution data are strongly correlated with those on corrupted data.

**Limitations:** While our experiments provide insight into souping, they are limited in scope. We only consider one dataset (CIFAR-100), one architecture (ResNet-50) with one baseline training trajectory. It is possible that our findings do not generalise to other datasets or architectures. Additionally, we only consider pairwise souping using the arithmetic mean at the midpoint. Other methods of souping, such as learned soups, may yield different results.

**Future Work:** Future empirical work could conduct similar experiments in different settings, such as a variety of model architectures and datasets. Theory could be developed for souping in simpler settings like an overparameterized linear model or a shallow network. Theory could also be created

216 to help characterise the variance reduction and consequently the reduction in loss we expect from  
 217 souping.  
 218

219 **REFERENCES**  
 220

221 Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models  
 222 modulo permutation symmetries, 2023. URL <https://arxiv.org/abs/2209.04836>.  
 223

224 Francesco Croce, Sylvestre-Alvise Rebuffi, Evan Shelhamer, and Sven Gowal. Seasoning model  
 225 soups for robustness to adversarial and natural distribution shifts, 2023. URL <https://arxiv.org/abs/2302.10164>.  
 226

227 Eduardo Dadalto. Resnet-50 model trained on cifar-100. [https://huggingface.co/edadaltocg/resnet50\\_cifar100](https://huggingface.co/edadaltocg/resnet50_cifar100), 2023. Hugging Face Model Repository.  
 228

229 Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. Linear mode  
 230 connectivity and the lottery ticket hypothesis, 2020. URL <https://arxiv.org/abs/1912.05671>.  
 231

232 Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep Residual Learning for Image Recog-  
 233 nition. In *Proceedings of 2016 IEEE Conference on Computer Vision and Pattern Recognition*,  
 234 CVPR '16, pp. 770–778, Las Vegas, NV, USA, June 2016. IEEE. doi: 10.1109/CVPR.2016.90.  
 235 URL <http://ieeexplore.ieee.org/document/7780459>.  
 236

237 Dan Hendrycks and Thomas Dietterich. Benchmarking neural network robustness to common  
 238 corruptions and perturbations, 2019. URL <https://arxiv.org/abs/1903.12261>.  
 239

240 Pavel Izmailov, Dmitrii Podoprikin, Timur Garipov, Dmitry Vetrov, and Andrew Gordon Wilson.  
 241 Averaging weights leads to wider optima and better generalization, 2019. URL <https://arxiv.org/abs/1803.05407>.  
 242

243 Dong-Hwan Jang, Sangdoo Yun, and Dongyoon Han. Model stock: All we need is just a few  
 244 fine-tuned models, 2025. URL <https://arxiv.org/abs/2403.19522>.  
 245

246 Daniel Morales-Brottons, Thijs Vogels, and Hadrien Hendrikx. Exponential moving average of  
 247 weights in deep learning: Dynamics and benefits, 2024. URL <https://arxiv.org/abs/2411.18704>.  
 248

249 Alexandre Ramé, Kartik Ahuja, Jianyu Zhang, Matthieu Cord, Léon Bottou, and David Lopez-Paz.  
 250 Model ratatouille: Recycling diverse models for out-of-distribution generalization, 2023. URL  
 251 <https://arxiv.org/abs/2212.10445>.  
 252

253 Mitchell Wortsman, Gabriel Ilharco, Samir Yitzhak Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes,  
 254 Ari S. Morcos, Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig  
 255 Schmidt. Model soups: averaging weights of multiple fine-tuned models improves accuracy  
 256 without increasing inference time, 2022. URL <https://arxiv.org/abs/2203.05482>.  
 257

258

259

260

261

262

263

264

265

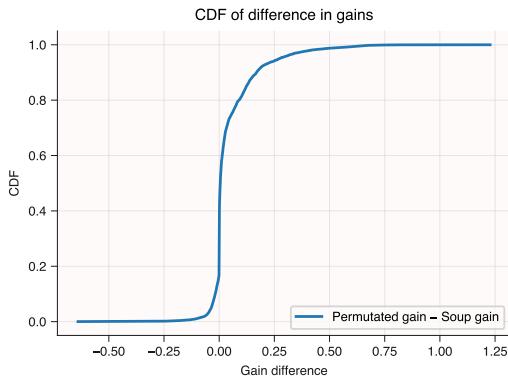
266

267

268

269

270  
271  
272  
273  
274  
275  
276  
277  
278  
279  
280  
281  
282



283  
284 Figure A1: Cumulative distribution function (CDF) of the difference in soup gain before and after  
285 permutation alignment using `rebasin`. This ignores the 7% of soups with a loss higher than 5 after  
286 permutation as these make the plot difficult to interpret. The remaining mean and median difference  
287 is approximately zero, indicating that permutation alignment does not have a significant effect on the  
288 effectiveness of souping in our experiments. While some soups benefit from permutation alignment,  
289 others are negatively affected, leading to an overall negligible impact while there is a risk of severe  
290 degradation.

## A APPENDIX

### A.1 PERMUTATION ALIGNMENT FOR SOUPING

295 Following the work from Ainsworth et al. (2023), we investigate whether permuting the neurons  
296 of models prior to souping increases the effectiveness of souping. We use the `rebasin`<sup>1</sup> package  
297 to align pairs of models before souping. This package uses the ‘matching weights’ method which  
298 permutes the neurones by inspecting only the weights. This contrasts with ‘activation matching’  
299 which requires forward passes through the network, and ‘straight through estimators’ which are  
300 even more computationally expensive. The authors find that matching weights performs similarly to  
301 activation matching while being computationally cheaper. Therefore, we only consider the matching  
302 weights method.

303 We align all 5,346 pairs of models using `rebasin` and compute the soup gain after alignment.  
304 Prior to permutation, 14.25% of soups were positive, while after permutation, 14.32% of soups were  
305 positive. However, 7% of soups obtained a loss higher than 5, which is worse than the loss of any  
306 previous soup. We plot the cumulative distribution function (CDF) of the difference in soup gain  
307 before and after permutation in Figure A1. This plot shows that while permuting can sometimes  
308 help, it does not do so consistently. Further, the median difference in soup gain is approximately  
309 zero, indicating that permuting does not have a significant effect on the effectiveness of souping in  
310 our experiments. There also remains a significant risk of severe degradation. Thus, we conclude that  
311 ‘matching weights’ permutation does not make a noticeable different to soupability in our setting.

312  
313  
314  
315  
316  
317  
318  
319  
320  
321  
322  
323

<sup>1</sup><https://pypi.org/project/rebasin/>

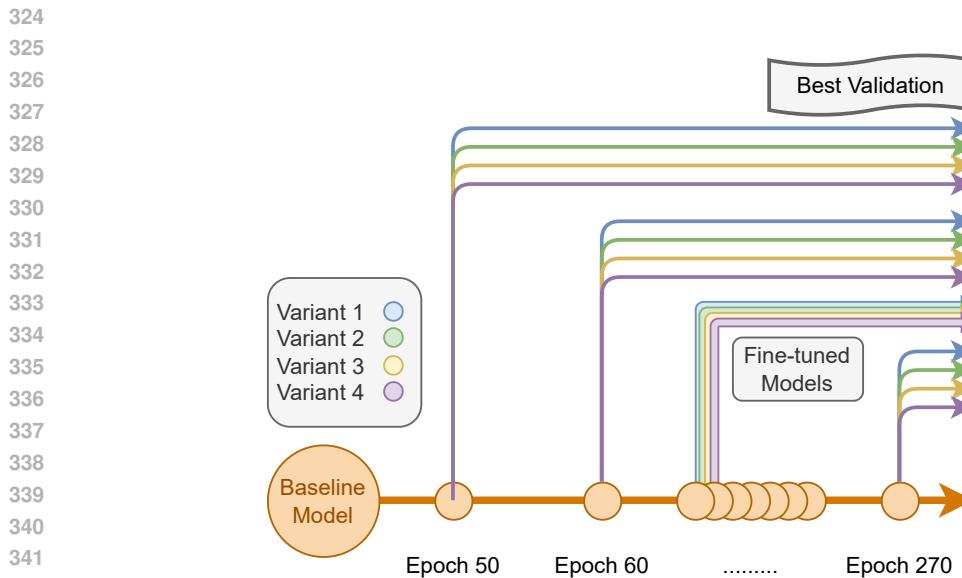


Figure A2: Branching of fine-tuned models from baseline checkpoints. A single baseline model is trained, with checkpoints saved every 10 epochs. From each checkpoint, 4 variants are trained with different optimizer hyper-parameter perturbations.

Model	Learning Rate Scale	Momentum Scale	Weight Decay Scale
Model 1	0.7073	1.1009	0.8284
Model 2	1.2244	0.9247	1.1150
Model 3	0.5112	1.0695	1.1099
Model 4	0.5373	0.8594	0.9078

Table 1: Optimizer perturbation scales applied during finetuning from the baseline ResNet-50 checkpoint on CIFAR-100. Each model scales the original SGD hyperparameters multiplicatively.

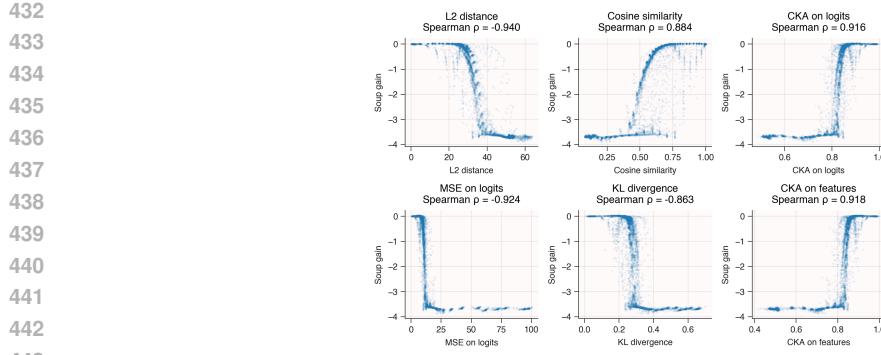
## A.2 TRAINING DETAILS FOR CIFAR-100 WITH RESNET-50

378  
 379  
 380  
 381  
 382  
 383  
 384  
 385  
 386  
 387  
 388  
 389  
 390  
 391  
 392  
 393  
 394

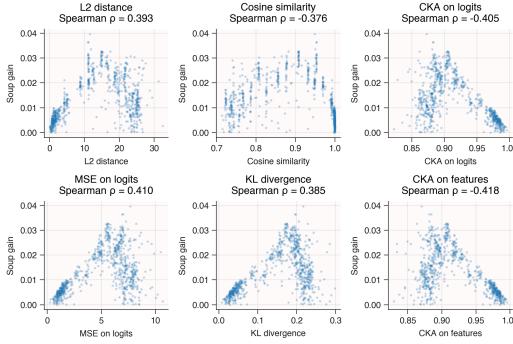
Component	Hyperparameter	Value
Dataset	Dataset	CIFAR-100
	# Classes	100
	Data augmentation	Mirroring and Padded Offset
	Validation split	5% of training set
	Split seed	42
Model	Architecture	ResNet-50
	Pretrained	No (from scratch)
Optimization	Optimizer	SGD (Nesterov)
	Initial learning rate	0.1
	Momentum	0.9
	Weight decay	$5 \times 10^{-4}$
Learning rate schedule	Scheduler	CosineAnnealingLR
	$T_{\max}$	280 epochs
	$\eta_{\min}$	0
Training	Epochs	300
	Batch size	128
	Mixed precision	No

Table 2: Training hyperparameters for CIFAR-100 with ResNet-50.

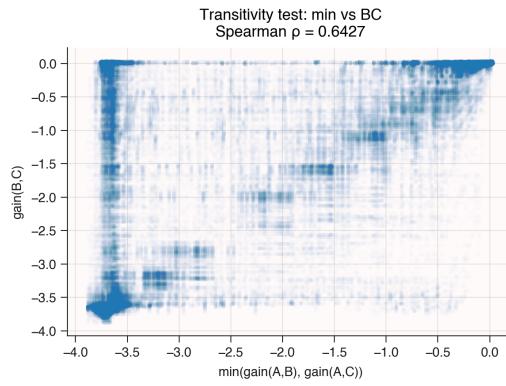
414  
 415  
 416  
 417  
 418  
 419  
 420  
 421  
 422  
 423  
 424  
 425  
 426  
 427  
 428  
 429  
 430  
 431



444 Figure A3: Soup gain vs various similarity and distance metrics between pairs of models. Each  
445 subplot shows the soup gain against one metric, with the Spearman correlation. All metrics perform  
446 similarly. We conclude that the more similar models are, the more likely souping is to not cause  
447 model collapse.



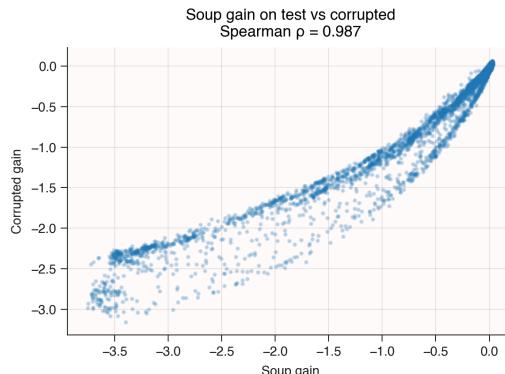
459 Figure A4: Soup gain vs various similarity and distance metrics between pairs of models, for the  
460 subset of soups with positive soup gain. Each subplot shows the soup gain against one metric, with  
461 the Spearman correlation. Each metric correlates similarly with soup gain. We see that when models  
462 are very similar, soup gain is small, while more dissimilar models can yield larger soup gains.  
463



479 Figure A5: Scatterplot of the soup gain of models  $B$  and  $C$  against the minimum soup gain of  
480 models  $A$  with  $B$  and  $C$ . Each point represents a triplet of models  $(A, B, C)$ . We observe a positive  
481 Spearman correlation of 0.64, suggesting that souping is fuzzily transitive. The correlation with the  
482 mean soup gain is lower, at 0.49.

483  
484  
485

486  
487  
488  
489  
490  
491  
492  
493  
494  
495  
496  
497  
498



499  
500  
501  
502  
503  
504  
505  
506  
507  
508  
509  
510  
511  
512  
513  
514  
515  
516  
517  
518  
519  
520  
521  
522  
523  
524  
525  
526  
527  
528  
529  
530  
531  
532  
533  
534  
535  
536  
537  
538  
539

Figure A6: Scatterplot of the soup gain on test vs corrupted data. There is a clear sub-linear trend with strong correlation. Such a close relationship is sensitive to the nature of the distribution shift. This plot mostly shows that when model collapse occurs on the original test set, it also occurs on the corrupted data.

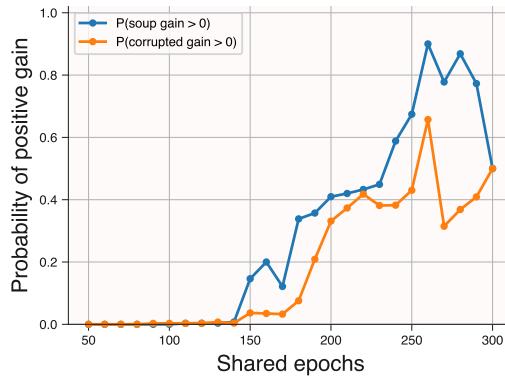


Figure A7: Probability of positive gain for soups as a function of the number of shared epochs. We see that the probability of positive gain increases with the number of shared epochs, for both clean and corrupted data. However, the corrupted data consistently has a slightly lower probability of positive gain. Thus, while souping also helps on corrupted data, it is slightly less effective than on clean data.

### A.3 FURTHER DETAILS ON EXPERIMENTS

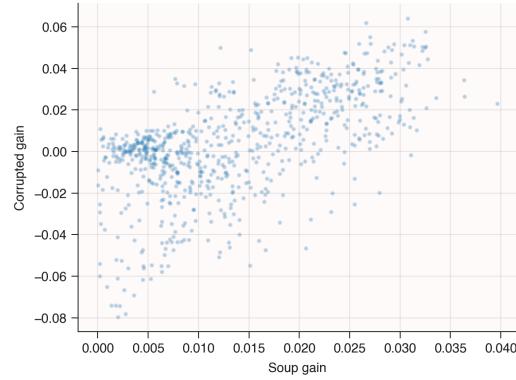
### A.4 DETAILS FOR MODEL EMBEDDING CREATION

Given our evidence for transitivity, we consider the broader landscape of all 104 of our originally trained models. Do the soups all exist in separate clusters of models in separate loss basins? We define a distance metric defined as

$$d_{AB} = -\text{sign}(\text{soup gain}) - 0.1 * \text{soup gain}$$

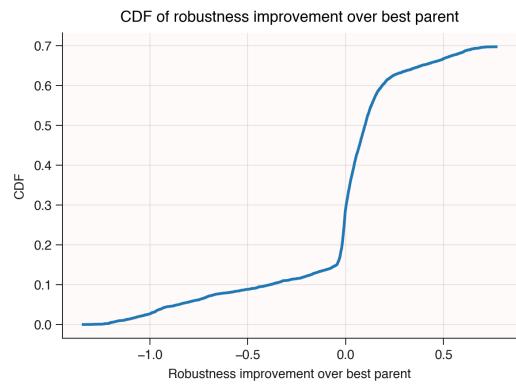
where  $d_{AB}$  is the distance between models  $A$  and  $B$ . Intuitively, this metric puts models close together that soup together positively, taking into account the magnitude of soup gain. We then cast this down into a 2-dimensional embedding using Multidimensional Scaling. We also color by branching epoch and mark edges that represent successful soups. The resulting plot is shown in Figure 6.

540  
541  
542  
543  
544  
545  
546  
547  
548  
549  
550  
551  
552  
553  
554  
555  
556  
557



558  
559 Figure A8: Plot of soup gain on test vs corrupted data for only models with positive soup gain on the  
560 test set. Spearman correlation of 0.61.

561  
562  
563  
564  
565  
566  
567  
568  
569  
570  
571  
572  
573  
574  
575  
576  
577  
578  
579  
580  
581  
582  
583



584  
585 Figure A9: CDF of the difference in *robustness gap* before and after souping. The robustness gap is  
586 defined as the difference in loss between the test set and the corrupted set. The robustness gap before  
587 souping is taken as the minimum robustness gap of the parents. We see a fairly symmetric distribution.  
588 It has mean -0.04 and median 0.02. We therefore conclude that souping does not systematically  
589 improve the robustness gap.  
590  
591  
592  
593