

WHEN IS MODEL SOUPING TASTY? SIMILARITY, TRANSITIVITY, AND ROBUSTNESS

000
001
002
003
004
005 **Anonymous authors**
006 Paper under double-blind review
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053
Model souping is a technique in which the parameters of models are averaged,
often leading to improved performance over constituent models without increasing
inference cost. However, the specific conditions required for success are not
well understood, particularly regarding the trade-off between model diversity and
stability. We analyse over 5,000 binary ResNet-50 soups trained on CIFAR-100,
with diversity controlled by branching ingredients from a shared training trajectory
at varying epochs. We find that effective souping requires a balance: models
must be similar enough to avoid model collapse, but diverse enough to yield
improvements. Furthermore, we provide empirical evidence for the hypothesis that
souping works by averaging within a low-loss basin. We also observe that soup
gains on corrupted data are strongly correlated with those on in-distribution data.
Finally, we compare souping to other parameter averaging methods and find that
XXX. Code and experiments are available at: <https://anonymous.4open.science/r/too-salty-478E/>.

ABSTRACT

1 INTRODUCTION

Wortsman et al. (2022) introduced the idea of *souping* by showing that averaging model parameters produced by different fine-tuning trajectories often leads to better generalization than any individual ingredient. This not only captures the benefit of multiple adaptation paths but also introduces an additional adaptive mechanism, as Croce et al. (2023) suggest that dynamically adjusting soup weights enables intermediate behaviours that can better match a range of distributional shifts. In contrast to ensembling which combines the *outputs* of models, souping does not increase the computational cost of inference, requiring only a single forward pass.

Wortsman et al. (2022) hypothesise that souping works because fine-tuned models often lie in the same low-loss basin. The convex combination of their weights is expected to remain in the low-loss basin while reducing variance introduced by a noisy training procedure. They find that when the angle formed by the pre-trained model and the two models to be souped is larger, there is a greater performance boost from souping. A wider angle suggests that fine-tuned trajectories are more diverse, thus more variance is reduced by averaging. Understanding why such averaging remains in low-loss regions, how diversity among fine-tuned models contributes to robustness, and when souping is beneficial is therefore essential for studying adaptation more broadly.

1.1 RELATED WORK

Souping: Souping has been used in a variety of settings. Croce et al. (2023) soup models trained to be robust to different distribution shifts and Ramé et al. (2023) soup ingredients trained on different tasks. Both cases lead to better generalisation. Jang et al. (2025) further explore the optimal soup between just two fine-tuned models by considering the angle formed by the two training trajectories in a layer-wise fashion.

Stochastic Weight Averaging (SWA): Souping is related to the idea of SWA Izmailov et al. (2019). In SWA, the ingredients of the soup come from different steps along the same training trajectory. By contrast, souping averages models from independent training trajectories. **Exponential Moving Average (EMA):** Another method for averaging weights over training is EMA Morales-Brotóns et al.

(2024). Here, the weights at each training step are combined with the previous average using an exponential decay.

Stability Analysis: Souping success requires models to be ‘compatible’ in the sense that averaging their weights has low loss. Frankle et al. (2020) define *stability to SGD noise*, whereby at some point during training, models become robust to the noise from SGD in the sense that all possible minima obtained by training from that point onwards lie in the same low-loss basin.

1.2 OUR CONTRIBUTIONS

Following these works, we seek to better understand souping by conducting a series of experiments addressing the following questions:

How much shared training is required for souping to be effective? We investigate varying the number of shared pre-training epochs before splitting into fine-tuned variants.

Can we predict the effectiveness of souping using similarity measures? If two models are very similar, it may be more likely that they can soup effectively.

Is souping transitive — If model A soups with B , and B soups with C , will A soup with C ?

Does souping in-distribution predict souping out-of-distribution? While souping has been shown to help with robustness to distribution shifts, we seek to answer how correlated the soup gains are between in-distribution and out-of-distribution data.

Further experiments investigating the effect of permuting models Ainsworth et al. (2023) prior to souping and how souping affects robustness can be found in Appendix A.1 and A9 respectively.

1.3 SOUPS AND SOUP GAIN

For this work, we only soup pairs of models θ_A and θ_B using the simple arithmetic mean. That is, $\theta_{\text{soup}} = \frac{1}{2}\theta_A + \frac{1}{2}\theta_B$. We do not consider any other weighted average in order to save on computational cost. Work by Ainsworth et al. (2023), shows many of the loss barriers they find have the most extreme behaviour at the midpoint. Thus we assume that the midpoint serves as an accurate summary statistic of souping performance over all possible convex combinations.

We define the *soup gain* of a pair of models θ_A, θ_B as the decrease in loss over a test set obtained by souping the models. That is, soup gain = $\min\{L(\theta_A), L(\theta_B)\} - L(\theta_{\text{soup}})$ where $L(\theta)$ denotes the test loss obtained using model θ . Soup gain can also be computed in terms of accuracy rather than loss. We use the soup gain as the primary measure of the effectiveness of souping in our experiments. We prefer to compare to the minimum rather than the mean of the ingredients as the purpose of a soup should be to improve over its ingredients.

2 EXPERIMENTS

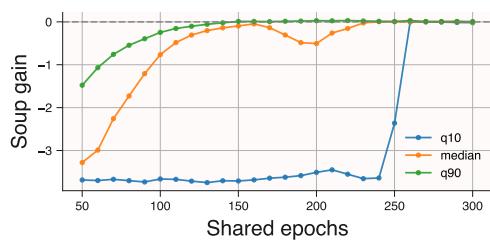
2.1 METHOD

We train a baseline model for image classification on the CIFAR-100 dataset with ResNet-50 He et al. (2016) using a well-known set of hyper-parameters Dadalto (2023) with image reflection and random translation with padding for data augmentation. We hold out 5% of the 50,000 training images as a validation set, as well as 10,000 images for testing. The baseline model is saved every 10 epochs. From each save point, we branch off and train 4 new models with different optimizer settings. For details, see Appendix Tables 1 and 2. All models are each trained to convergence, with the best validation scored model weights saved for the experiments. This process is illustrated in Figure A2. A total of 4 variants and 26 branch points were trained, yielding 104 related models and 5,356 binary souping combinations for analysis.

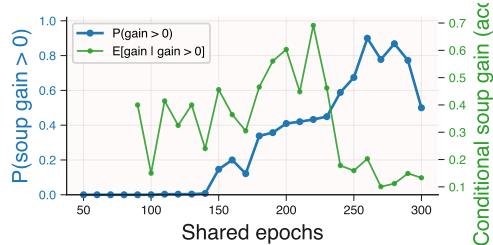
2.2 SHARED EPOCHS AND SOUP GAIN

We group the soups by the number of shared epochs between the two ingredients before they diverge into different training trajectories. For example, a pair of models branched form the baseline at epoch 50 and 100 respectively share 50 epochs. In Figure 1, we plot quantiles of soup gain as a function of shared epochs. We observe the distribution of soup gains shifts positively as the number of shared

108 epochs increases. However, the soup gain is often large and negative until around epoch 150. After
 109 epoch 250, nearly all soups are approximately neutral.
 110



120 Figure 1: Quantiles of soup gain vs shared
 121 epochs.

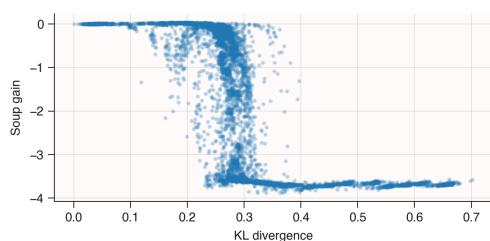


122 Figure 2: Probability of positive soup gain and
 123 conditional expected gain vs shared epochs.

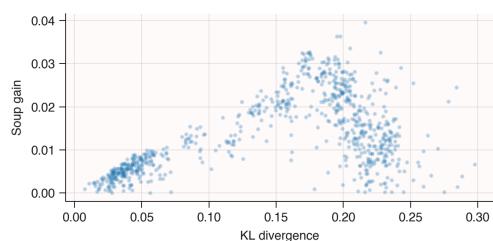
124 To better illustrate how much souping helps, we plot the probability of positive soup gain and the
 125 average soup gain in accuracy for such positive soups in Figure 2. The probability of positive soup
 126 gain increases with shared epochs, reaching around 80% after 250 shared epochs. The conditional
 127 expected soup gain is noisy, at around 0.5% accuracy improvement when souping works. Towards
 128 the upper end of shared epochs, the gains decrease to around 0.2%. This suggests that while souping
 129 becomes more likely to work with increased shared training, the magnitude of gain decreases when
 130 the models are too similar.

131 2.3 PREDICTING SOUPABILITY WITH SIMILARITY

132 Given two models, can we predict whether or not they will soup? To test this hypothesis, we compute
 133 a variety of similarity and distance metrics between pairs of models. We find that all metrics perform
 134 similarly. A plot for all metrics can be found in Figure A3. We arbitrarily choose to show the KL
 135 divergence between the outputs of the ingredients in Figure 3 as an example. There is a strong
 136 correlation between KL divergence and soup gain with a Spearman correlation of -0.86.
 137 However, many of these soups have very poor performance. Figure 4 shows the soup gain against KL
 138 divergence for only those models with positive soup gain, with a moderate positive correlation with a
 139 Spearman correlation of 0.39. Models must be sufficiently similar in order to soup, but to be effective,
 140 the models must also be sufficiently different. Balancing these two effects is key to tasty soups.



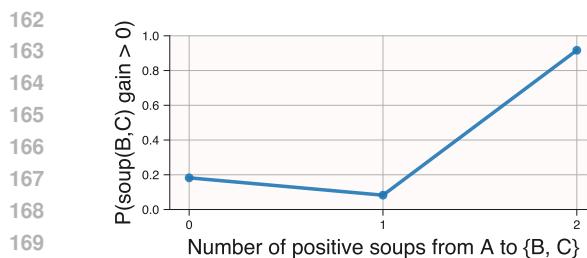
141 Figure 3: KL vs soup gain (Spearman -0.86).



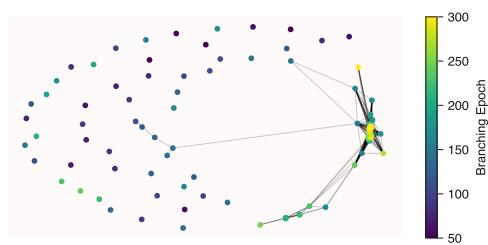
142 Figure 4: KL vs soup gain, positive soups only
 143 (Spearman 0.39).

144 2.4 IS SOUPING TRANSITIVE?

145 If model A soups successfully with models B and C , will B and C also soup successfully? If models
 146 soups together when they lie in the same low loss region, then transitivity should hold. To test this
 147 hypothesis, we consider all triplets of models (A, B, C) . We plot the probability that B and C soup
 148 against the number of positive soups involving A in Figure 5. We observe that the probability that B
 149 and C soup is only high when A soups with both B and C . We also plot the soup gain between B
 150 and C against the minimum soup gain of (A, B) and (A, C) in Figure A5, noting a moderate positive
 151 Spearman correlation of 0.64.



171 Figure 5: Probability of positive soup gain of B
172 and C vs positive soups with A .
173



171 Figure 6: 2D embedding of 104 models using
172 soup-gain distance; edges indicate positive
173 soups.
174
175

Following transitivity, we investigate whether there are clusters of models that soup well together. We embed our 104 models into 2D using soup gain as a distance metric and plot them in Figure 6. The details of generating this plot can be found in Section A.4. We find that most successful soups lie in a single cluster of models all branched from later on in the training procedure. We conclude that souping is approximately transitive. There are many counter-examples, but this work supports the idea that, in general, soups lie in a single loss basin.

2.5 SOUPING FOR ROBUSTNESS TO CORRUPTION

All experiments thus far have measured soup gain on a held-out test set. However, souping has also been used for robustness to distribution shift (Croce et al., 2023). To establish whether souping for in-distribution performance also increases out-of-distribution performance, we compute the soup gain on CIFAR-100C (Hendrycks & Dietterich, 2019) with severity level 3. The soup gains on test and corrupted data correlate very well, with a Spearman correlation of 0.99. A scatterplot can be found in Figure A6. The positive trend still holds when conditioning on only soups with positive soup gain on the test set, with a Pearson correlation of 0.61. A plot can be found in Figure A8. In-distribution performance improvement transfers to unseen target distributions.

We also plot the probability of positive soup gain on corrupted data as a function of shared epochs in Figure A7. The probability of positive gain increases with the number of shared epochs for both clean and corrupted data but the corrupted data always has a lower probability of positive gain.

3 CONCLUSION

In this work, we have conducted a series of experiments to further our understanding of souping. We find that ingredients must have sufficiently many shared epochs of training in order to be compatible but that too many shared epochs and the soup gain is minimal. We also find that various similarity measures between models correlate similarly with soup gain. Similar ingredients are less likely to collapse when souped, but very similar ingredients yield smaller soup gains. Thus, the right balance must be struck for the most effective souping. When souping, we encourage practitioners to test a range of similarities of ingredients to ensure they are finding an optimal soup. Our experiments showing that souping is mostly transitive support the low-loss basin hypothesis. Finally, we find that soup gains on in-distribution data are strongly correlated with those on corrupted data.

Limitations: While our experiments provide insight into souping, they are limited in scope. We only consider one dataset (CIFAR-100), one architecture (ResNet-50) with one baseline training trajectory. It is possible that our findings do not generalise to other datasets or architectures. Additionally, we only consider pairwise souping using the arithmetic mean at the midpoint. Other methods of souping, such as learned soups, may yield different results.

Future Work: Future empirical work could conduct similar experiments in different settings, such as a variety of model architectures and datasets. Theory could be developed for souping in simpler settings like an overparameterized linear model or a shallow network. Theory could also be created to help predict the change in loss we expect from souping associated with noise reduction.

216 REFERENCES
217218 Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models
219 modulo permutation symmetries, 2023. URL <https://arxiv.org/abs/2209.04836>.220 Francesco Croce, Sylvestre-Alvise Rebuffi, Evan Shelhamer, and Sven Gowal. Seasoning model
221 soups for robustness to adversarial and natural distribution shifts, 2023. URL <https://arxiv.org/abs/2302.10164>.222 Eduardo Dadalto. Resnet-50 model trained on cifar-100. https://huggingface.co/edadaltocg/resnet50_cifar100, 2023. Hugging Face Model Repository.223 Jonathan Frankle, Gintare Karolina Dziugaite, Daniel M. Roy, and Michael Carbin. Linear mode
224 connectivity and the lottery ticket hypothesis, 2020. URL <https://arxiv.org/abs/1912.05671>.225 Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep Residual Learning for Image Recog-
226 nition. In *Proceedings of 2016 IEEE Conference on Computer Vision and Pattern Recognition*,
227 CVPR '16, pp. 770–778, Las Vegas, NV, USA, June 2016. IEEE. doi: 10.1109/CVPR.2016.90.
228 URL <http://ieeexplore.ieee.org/document/7780459>.229 Dan Hendrycks and Thomas Dietterich. Benchmarking neural network robustness to common
230 corruptions and perturbations, 2019. URL <https://arxiv.org/abs/1903.12261>.231 Pavel Izmailov, Dmitrii Podoprikhin, Timur Garipov, Dmitry Vetrov, and Andrew Gordon Wilson.
232 Averaging weights leads to wider optima and better generalization, 2019. URL <https://arxiv.org/abs/1803.05407>.233 Dong-Hwan Jang, Sangdoo Yun, and Dongyoon Han. Model stock: All we need is just a few
234 fine-tuned models, 2025. URL <https://arxiv.org/abs/2403.19522>.235 Daniel Morales-Brotos, Thijs Vogels, and Hadrien Hendrikx. Exponential moving average of
236 weights in deep learning: Dynamics and benefits, 2024. URL <https://arxiv.org/abs/2411.18704>.237 Alexandre Ramé, Kartik Ahuja, Jianyu Zhang, Matthieu Cord, Léon Bottou, and David Lopez-Paz.
238 Model ratatouille: Recycling diverse models for out-of-distribution generalization, 2023. URL
239 <https://arxiv.org/abs/2212.10445>.240 Mitchell Wortsman, Gabriel Ilharco, Samir Yitzhak Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes,
241 Ari S. Morcos, Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig
242 Schmidt. Model soups: averaging weights of multiple fine-tuned models improves accuracy
243 without increasing inference time, 2022. URL <https://arxiv.org/abs/2203.05482>.

244

245

246

247

248

249

250

251

252

253

254

255

256

257

258

259

260

261

262

263

264

265

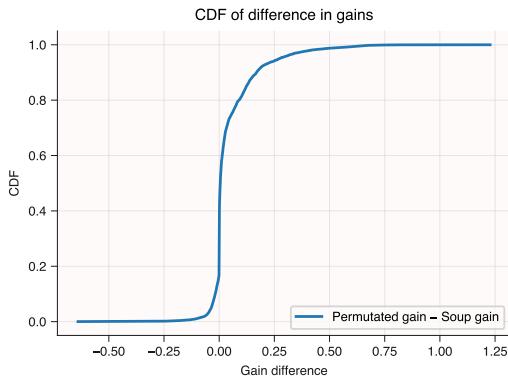
266

267

268

269

270
271
272
273
274
275
276
277
278
279
280
281
282



283
284 Figure A1: Cumulative distribution function (CDF) of the difference in soup gain before and after
285 permutation alignment using `rebasin`. This ignores the 7% of soups with a loss higher than 5 after
286 permutation as these make the plot difficult to interpret. The remaining mean and median difference
287 is approximately zero, indicating that permutation alignment does not have a significant effect on the
288 effectiveness of souping in our experiments. While some soups benefit from permutation alignment,
289 others are negatively affected, leading to an overall negligible impact while there is a risk of severe
290 degradation.

A APPENDIX

A.1 PERMUTATION ALIGNMENT FOR SOUPING

295 Following the work from Ainsworth et al. (2023), we investigate whether permuting the neurons
296 of models prior to souping increases the effectiveness of souping. We use the `rebasin`¹ package
297 to align pairs of models before souping. This package uses the ‘matching weights’ method which
298 permutes the neurones by inspecting only the weights. This contrasts with ‘activation matching’
299 which requires forward passes through the network, and ‘straight through estimators’ which are
300 even more computationally expensive. The authors find that matching weights performs similarly to
301 activation matching while being computationally cheaper. Therefore, we only consider the matching
302 weights method.

303 We align all 5,346 pairs of models using `rebasin` and compute the soup gain after alignment.
304 Prior to permutation, 14.25% of soups were positive, while after permutation, 14.32% of soups were
305 positive. However, 7% of soups obtained a loss higher than 5, which is worse than the loss of any
306 previous soup. We plot the cumulative distribution function (CDF) of the difference in soup gain
307 before and after permutation in Figure A1. This plot shows that while permuting can sometimes
308 help, it does not do so consistently. Further, the median difference in soup gain is approximately
309 zero, indicating that permuting does not have a significant effect on the effectiveness of souping in
310 our experiments. There also remains a significant risk of severe degradation. Thus, we conclude that
311 ‘matching weights’ permutation does not make a noticeable different to soupability in our setting.

312
313
314
315
316
317
318
319
320
321
322
323

¹<https://pypi.org/project/rebasin/>

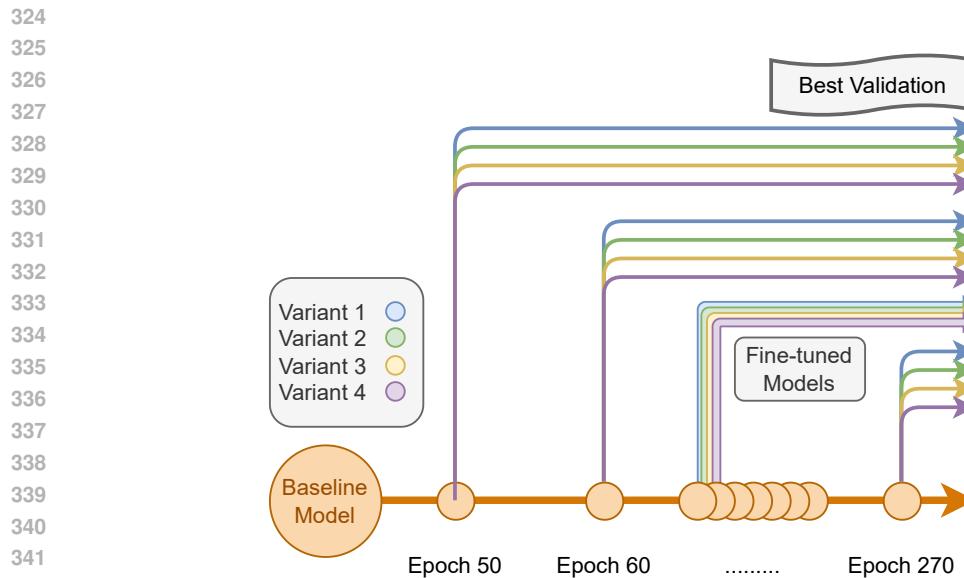


Figure A2: Branching of fine-tuned models from baseline checkpoints. A single baseline model is trained, with checkpoints saved every 10 epochs. From each checkpoint, 4 variants are trained with different optimizer hyper-parameter perturbations.

Model	Learning Rate Scale	Momentum Scale	Weight Decay Scale
Model 1	0.7073	1.1009	0.8284
Model 2	1.2244	0.9247	1.1150
Model 3	0.5112	1.0695	1.1099
Model 4	0.5373	0.8594	0.9078

Table 1: Optimizer perturbation scales applied during finetuning from the baseline ResNet-50 checkpoint on CIFAR-100. Each model scales the original SGD hyperparameters multiplicatively.

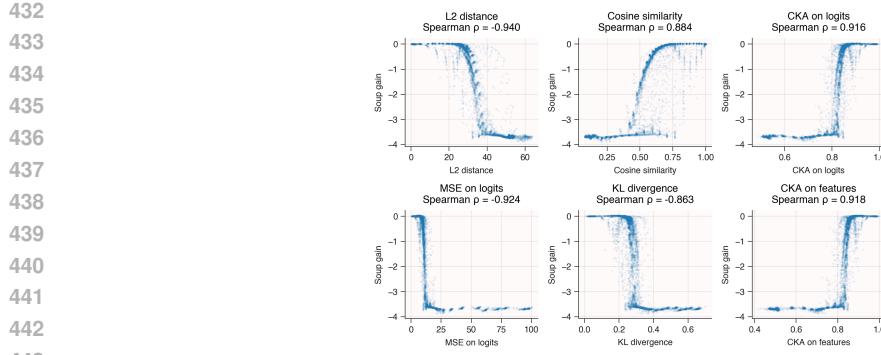
A.2 TRAINING DETAILS FOR CIFAR-100 WITH RESNET-50

378
 379
 380
 381
 382
 383
 384
 385
 386
 387
 388
 389
 390
 391
 392
 393
 394


Component	Hyperparameter	Value
Dataset	Dataset	CIFAR-100
	# Classes	100
	Data augmentation	Mirroring and Padded Offset
	Validation split	5% of training set
	Split seed	42
Model	Architecture	ResNet-50
	Pretrained	No (from scratch)
Optimization	Optimizer	SGD (Nesterov)
	Initial learning rate	0.1
	Momentum	0.9
	Weight decay	5×10^{-4}
Learning rate schedule	Scheduler	CosineAnnealingLR
	T_{\max}	280 epochs
	η_{\min}	0
Training	Epochs	300
	Batch size	128
	Mixed precision	No

Table 2: Training hyperparameters for CIFAR-100 with ResNet-50.


414
 415
 416
 417
 418
 419
 420
 421
 422
 423
 424
 425
 426
 427
 428
 429
 430
 431



444 Figure A3: Soup gain vs various similarity and distance metrics between pairs of models. Each
445 subplot shows the soup gain against one metric, with the Spearman correlation. All metrics perform
446 similarly. We conclude that the more similar models are, the more likely souping is to not cause
447 model collapse.



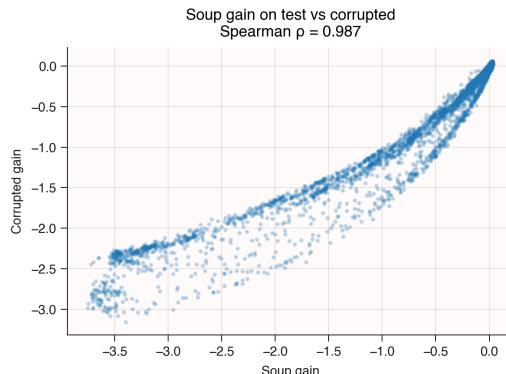
460 Figure A4: Soup gain vs various similarity and distance metrics between pairs of models, for the
461 subset of soups with positive soup gain. Each subplot shows the soup gain against one metric, with
462 the Spearman correlation. Each metric correlates similarly with soup gain. We see that when models
463 are very similar, soup gain is small, while more dissimilar models can yield larger soup gains.



479 Figure A5: Scatterplot of the soup gain of models B and C against the minimum soup gain of
480 models A with B and C . Each point represents a triplet of models (A, B, C) . We observe a positive
481 Spearman correlation of 0.64, suggesting that souping is fuzzily transitive. The correlation with the
482 mean soup gain is lower, at 0.49.

483
484
485

486
487
488
489
490
491
492
493
494
495
496
497
498



499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539

Figure A6: Scatterplot of the soup gain on test vs corrupted data. There is a clear sub-linear trend with strong correlation. Such a close relationship is sensitive to the nature of the distribution shift. This plot mostly shows that when model collapse occurs on the original test set, it also occurs on the corrupted data.

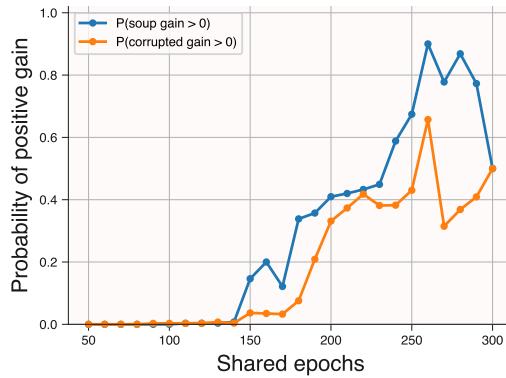


Figure A7: Probability of positive gain for soups as a function of the number of shared epochs. We see that the probability of positive gain increases with the number of shared epochs, for both clean and corrupted data. However, the corrupted data consistently has a slightly lower probability of positive gain. Thus, while souping also helps on corrupted data, it is slightly less effective than on clean data.

A.3 FURTHER DETAILS ON EXPERIMENTS

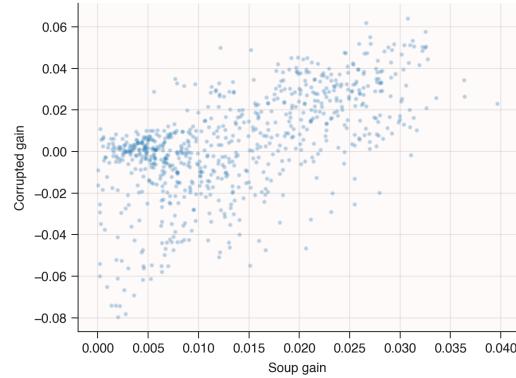
A.4 DETAILS FOR MODEL EMBEDDING CREATION

Given our evidence for transitivity, we consider the broader landscape of all 104 of our originally trained models. Do the soups all exist in separate clusters of models in separate loss basins? We define a distance metric defined as

$$d_{AB} = -\text{sign}(\text{soup gain}) - 0.1 * \text{soup gain}$$


where d_{AB} is the distance between models A and B . Intuitively, this metric puts models close together that soup together positively, taking into account the magnitude of soup gain. We then cast this down into a 2-dimensional embedding using Multidimensional Scaling. We also color by branching epoch and mark edges that represent successful soups. The resulting plot is shown in Figure 6.

540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557



558
559 Figure A8: Plot of soup gain on test vs corrupted data for only models with positive soup gain on the
560 test set. Spearman correlation of 0.61.

561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583



584
585 Figure A9: CDF of the difference in *robustness gap* before and after souping. The robustness gap is
586 defined as the difference in loss between the test set and the corrupted set. The robustness gap before
587 souping is taken as the minimum robustness gap of the parents. We see a fairly symmetric distribution.
588 It has mean -0.04 and median 0.02. We therefore conclude that souping does not systematically
589 improve the robustness gap.
590
591
592
593